

COMPSCI 589 – Final Project

Benchmarking Classical ML Algorithms on Four Datasets

Vishal Hanuman

Akshith Karthik

May 8, 2026

Note on Implementations

All machine learning algorithms were implemented from scratch. No machine learning libraries (e.g., scikit-learn, PyTorch) were used for training or prediction. `numpy` and `matplotlib` were used for numerical computation and plotting. `scikit-learn` was used *only* for dataset loading, not for training or prediction.

- **Datasets 1–2** (Digits, Parkinson’s): k-NN, Gaussian Naive Bayes, Random Forest, and Neural Network implementations by **Vishal Hanuman** (HW1–HW4).
- **Datasets 3–4** (Rice, Credit Approval): Experiments by **Akshith Karthik**. Rice uses k-NN, Decision Tree, and Gaussian Naive Bayes implementations from Akshith Karthik’s homework code. Credit Approval uses Akshith Karthik’s Decision Tree and Random Forest implementations, with the shared HW4 Neural Network used for the additional neural-network comparison.

1 Introduction

The goal of this project is to compare the performance of classical machine learning algorithms on four benchmark datasets by evaluating accuracy and F1-score under stratified 10-fold cross-validation, with systematic hyperparameter tuning.

2 Datasets

1. **Digits** – 1,797 instances, 64 numerical features, 10 classes (0–9). Loaded using the Digits loader from `scikit-learn`.
2. **Parkinson’s Disease** – 195 instances, 22 numerical features, binary (0 = healthy, 1 = Parkinson’s).
3. **Rice Grains** – 3,810 instances, 7 numerical features, binary (Cammeo vs. Osmancik).
4. **Credit Approval** – 653 instances, 6 numerical + 9 categorical features, binary (approved vs. not).

3 Experiments and Analyses

3.1 Dataset 1 – Hand-Written Digits

For the Digits dataset, we evaluated k-NN, Random Forest, and a Neural Network. All features are numerical pixel intensities, so no feature encoding was required. k-NN is a natural baseline because digits with similar pixel layouts should be close in the normalized feature space. Random Forests were tested because they can learn nonlinear decision rules over pixel thresholds, while the Neural Network was included because it can combine many pixel-level signals through hidden units. All results below use stratified 10-fold cross-validation and report accuracy and macro F1-score.

k-NN Hyperparameter Sweep

Table 1: Digits: k-NN performance for different values of k .

k	Accuracy	Macro F1
1	0.9872	0.9871
3	0.9872	0.9872
5	0.9883	0.9883
7	0.9861	0.9861
11	0.9794	0.9793
15	0.9806	0.9805
21	0.9745	0.9743
31	0.9677	0.9674

Random Forest Hyperparameter Sweep

Table 2: Digits: Random Forest performance for different values of $ntree$.


$ntree$	Accuracy	Macro F1
5	0.8965	0.8957
10	0.9377	0.9373
20	0.9605	0.9604
30	0.9621	0.9620
50	0.9694	0.9695
75	0.9716	0.9718
100	0.9660	0.9660
150	0.9749	0.9748

Neural Network Hyperparameter Sweep


Table 3: Digits: Neural Network performance for selected architectures and regularization values.

Hidden layers	λ	Accuracy	Macro F1
(32,)	0.0	0.9800	0.9798
(64,)	0.0	0.9800	0.9799
(32, 16)	0.0	0.9756	0.9755
(64, 32)	0.0	0.9761	0.9761
(32,)	0.1	0.9544	0.9542
(64, 32)	0.1	0.9538	0.9536

Learning Curves




(a) k-NN: Accuracy and F1 vs k




(b) Random Forest: Accuracy and F1 vs $ntree$

Figure 1: Hyperparameter curves for Dataset 1 (Digits).



(a) Neural Network: Accuracy and F1 by configuration



(b) Neural Network: Cost vs training set size

Figure 2: Neural Network graphs for Dataset 1 (Digits).

k-NN achieved the best Digits performance, with $k = 5$ producing 0.9883 accuracy and 0.9883 macro F1. Very small values of k already worked well, which suggests that the normalized pixel space has strong local class structure. Performance declined as k became larger, likely because broad neighborhoods smooth over class boundaries between visually similar digits.

Random Forest performance improved substantially as the number of trees increased, but it remained below k-NN and the Neural Network. This is reasonable for image-like data: an axis-aligned tree split only sees one pixel feature at a time, so the forest needs many trees to approximate

spatial digit patterns. The Neural Network performed strongly, with the best setting at one hidden layer of 64 units and no regularization. Adding regularization with $\lambda = 0.1$ reduced performance, which indicates that these small networks were more likely underfitting than overfitting in this experiment. The learning curve decreased as more training examples were used and then began to flatten, suggesting that the selected network was converging but had less advantage over the simpler k-NN classifier on this dataset.

3.2 Dataset 2 – Parkinson’s Disease

For the Parkinson’s dataset, we evaluated k-NN, Gaussian Naive Bayes, Random Forest, and a Neural Network. The dataset contains 22 numerical voice-measurement features and a binary target variable, **Diagnosis**. k-NN tests whether nearby voice measurements identify similar diagnoses. Naive Bayes provides a simple probabilistic baseline, Random Forests capture nonlinear feature thresholds, and the Neural Network tests whether a learned nonlinear representation improves performance. The F1-score reported here is the F1-score for the positive Parkinson’s class.

k-NN Hyperparameter Sweep

Table 4: Parkinson’s: k-NN performance for different values of k .

k	Accuracy	F1
1	0.9544	0.9676
3	0.9278	0.9530
5	0.9278	0.9528
7	0.9325	0.9565
11	0.8917	0.9319
15	0.8420	0.9034
21	0.8373	0.9028
31	0.8403	0.9047

Gaussian Naive Bayes Hyperparameter Sweep

Table 5: Parkinson’s: Gaussian Naive Bayes performance for different variance floors.

ϵ	Accuracy	F1
1e-9	0.6915	0.7423
1e-7	0.6968	0.7479
1e-5	0.6979	0.7504
0.001	0.6929	0.7497
0.01	0.7129	0.7651
0.1	0.7171	0.7695

Random Forest Hyperparameter Sweep

Table 6: Parkinson's: Random Forest performance for different values of $ntree$.


$ntree$	Accuracy	F1
5	0.8554	0.9055
10	0.8817	0.9245
20	0.9014	0.9376
30	0.8878	0.9286
50	0.8920	0.9327
75	0.8823	0.9250
100	0.8928	0.9324
150	0.8911	0.9311

Neural Network Hyperparameter Sweep


Table 7: Parkinson's: Neural Network performance for selected architectures and regularization values.

Hidden layers	λ	Accuracy	F1
(8,)	0.0	0.8323	0.8920
(16,)	0.0	0.8378	0.8968
(8, 4)	0.0	0.7648	0.8652
(16, 8)	0.0	0.7659	0.8627
(8,)	0.1	0.8303	0.8966
(16, 8)	0.1	0.7542	0.8598

Learning Curves




(a) k-NN: Accuracy and F1 vs k




(b) Naive Bayes: Accuracy and F1 vs ϵ

Figure 3: Hyperparameter curves for k-NN and Naive Bayes on Dataset 2.



(a) Random Forest: Accuracy and F1 vs *ntree*



(b) Neural Network: Accuracy and F1 by configuration

Figure 4: Random Forest and Neural Network configuration curves for Dataset 2.




Figure 5: Neural Network learning curve for Dataset 2.

k-NN achieved the best Parkinson's performance, with $k = 1$ producing 0.9544 accuracy and 0.9676 F1. Larger values of k reduced performance, which suggests that the most useful signal is very local in the normalized voice-feature space. Because the dataset is small and imbalanced toward the positive class, larger neighborhoods likely mix less similar patients and smooth away useful diagnostic differences.

Gaussian Naive Bayes was the weakest model. The voice measurements are likely correlated with each other, while Naive Bayes assumes conditional independence given the class. The epsilon sweep shows that increasing the variance floor helped slightly, but it did not close the gap to the other methods. Random Forest was the second-best method: it handled nonlinear feature thresholds better than Naive Bayes, but performance peaked at 20 trees and then varied slightly, which is consistent with the higher variance expected from bootstrap training on only 195 instances. The Neural Network trailed k-NN and Random Forest; the deeper architectures were worse, suggesting that the dataset is too small for the added capacity. Its learning curve generally decreased as more training instances were used, but the curve was not perfectly smooth because the train/test split is small and the training procedure is stochastic.

3.3 Dataset 3 – Rice Grains

For the Rice Grains dataset, we evaluated k-NN, a Decision Tree, and Gaussian Naive Bayes. The dataset has 3,810 examples with seven numerical shape descriptors and a binary label (*Cammeo* or *Osmancik*). Since all features are numerical, k-NN can use normalized distances directly. The Decision Tree provides a comparison based on learned feature-threshold rules, while Gaussian Naive

Bayes tests whether a simple class-conditional Gaussian model is sufficient for the rice shape descriptors. We used stratified 10-fold cross-validation and reported accuracy and F1-score.

k-NN Hyperparameter Sweep

Table 8: Rice: k-NN performance for different values of k .

k	Accuracy	F1
1	0.8861	0.8999
3	0.9113	0.9226
5	0.9226	0.9325
7	0.9231	0.9330
9	0.9244	0.9342
11	0.9239	0.9337

Decision Tree Hyperparameter Sweep

Table 9: Rice: Decision Tree performance for different maximum depths.

Max depth	Accuracy	F1
2	0.8958	0.9027
3	0.9121	0.9267
4	0.9278	0.9364
5	0.9257	0.9352
6	0.9294	0.9385
7	0.9273	0.9364

Gaussian Naive Bayes Hyperparameter Sweep

Table 10: Rice: Gaussian Naive Bayes performance for different variance floors.

ϵ	Accuracy	F1
1e-9	0.9171	0.9024
1e-7	0.9168	0.9021
1e-5	0.9176	0.9032
0.001	0.9171	0.9024
0.01	0.9171	0.9024
0.1	0.9105	0.8946

Learning Curves




Figure 6: Rice: k-NN, Decision Tree, and Gaussian Naive Bayes hyperparameter sweeps.

All three methods performed reasonably well on Rice, but the Decision Tree gave the best overall result. A maximum depth of 6 reached 0.9294 accuracy and 0.9385 F1, suggesting that a moderate number of threshold-based rules captures the separation between the two rice varieties better than a very shallow tree. k-NN was close behind: $k = 9$ produced 0.9244 accuracy and 0.9342 F1, and the results stayed stable from $k = 5$ through $k = 11$. This indicates that local neighborhoods are informative, but the dataset benefits from some smoothing instead of relying on a single nearest neighbor. Gaussian Naive Bayes was weaker than the tree and k-NN, with its best result at $\epsilon = 10^{-5}$. Its lower F1 suggests that the rice shape features are not conditionally independent, although the model still captured much of the class structure from the numerical descriptors.

3.4 Dataset 4 – Credit Approval

For the Credit Approval dataset, we evaluated Decision Tree, Random Forest, and Neural Network models. This dataset is different from the first three because it mixes numerical and categorical attributes. The categorical variables were converted into one-hot features before model training, while numerical variables were kept as continuous values. We again used stratified 10-fold cross-validation and compared accuracy and F1-score.

Decision Tree Hyperparameter Sweep

Table 11: Credit Approval: Decision Tree performance for different maximum depths.

Max depth	Accuracy	F1
2	0.8635	0.8627
4	0.8359	0.8263
6	0.8178	0.8059
8	0.8101	0.7975
10	0.8178	0.8006
12	0.7932	0.7729

Random Forest Hyperparameter Sweep

Table 12: Credit Approval: Random Forest performance for different values of *ntree*.

<i>ntree</i>	Accuracy	F1
1	0.8311	0.8191
5	0.8652	0.8517
10	0.8604	0.8507
15	0.8605	0.8492
20	0.8666	0.8552
25	0.8682	0.8583

Neural Network Hyperparameter Sweep

Table 13: Credit Approval: Neural Network performance for selected architectures and regularization values.

Hidden layers	λ	Accuracy	F1
(8,)	0.0	0.8544	0.8389
(16,)	0.0	0.8559	0.8415
(8, 4)	0.0	0.8559	0.8421
(16, 8)	0.0	0.8467	0.8268
(8,)	0.1	0.8665	0.8605
(16, 8)	0.1	0.5467	0.0000

Learning Curves






Figure 7: Credit Approval: Decision Tree and Random Forest hyperparameter sweeps.



(a) Neural Network: Accuracy and F1 by configuration



(b) Neural Network: Cost vs training set size

Figure 8: Credit Approval: Neural Network configuration and learning curves.

The Credit Approval results were tightly grouped among the best configurations. Random Forest gave the highest accuracy, with $ntree=25$ producing 0.8682 accuracy and 0.8583 F1. The single Decision Tree was very close and produced the highest F1-score: depth 2 reached 0.8635 accuracy and 0.8627 F1, but deeper trees declined. That pattern suggests that the dataset has a few strong approval rules, while deeper splits begin to memorize smaller subgroups. The Neural Network was also competitive. The best network used one hidden layer of 8 units with $\lambda = 0.1$, reaching 0.8665 accuracy and 0.8605 F1. The larger regularized architecture collapsed to an F1 of 0.0000, which likely means it predicted mostly one class under that setting. Overall, a shallow single tree and a moderately sized forest fit this mixed-feature dataset slightly better than the deeper neural-network configurations.

4 Summary of Results

Table 14: Best performance (Accuracy / F1-Score) per algorithm per dataset. Gray cells indicate the highest value for each metric on each dataset.

Algorithm	Digits		Parkinson's		Rice		Credit	
	Acc	F1	Acc	F1	Acc	F1	Acc	F1
k-NN	0.9883	0.9883	0.9544	0.9676	0.9244	0.9342	—	—
Decision Tree	—	—	—	—	0.9294	0.9385	0.8635	0.8627
Naive Bayes	—	—	0.7171	0.7695	0.9176	0.9032	—	—
Random Forest	0.9749	0.9748	0.9014	0.9376	—	—	0.8682	0.8583
Neural Network	0.9800	0.9799	0.8378	0.8968	—	—	0.8665	0.8605

5 Discussion

Across the four main datasets, no single model dominated every setting, but the patterns were consistent with the structure of the data. k-NN was strongest on Digits and Parkinson's, and it was also competitive on Rice. These datasets all have numerical feature spaces where distance is meaningful after normalization. Digits and Rice especially show that local neighborhoods carry strong class information, although the Rice results show that a moderately deep threshold tree can be even stronger.

Tree-based methods were strongest on Credit Approval. That dataset mixes categorical and numerical features, and the approval label appears to depend on a small set of rule-like conditions. A shallow Decision Tree already performed well, and the Random Forest improved slightly by averaging several bootstrapped trees. In contrast, deeper single trees became worse on Credit, which points to overfitting rather than underfitting. Rice was different: the tree improved through moderate depth and then leveled off, suggesting that several shape thresholds are useful before extra depth begins to add little.

The Neural Network was competitive when there were enough examples and a useful continuous feature representation, especially on Digits. It was weaker on Parkinson’s and Credit. For Parkinson’s, the dataset is very small, so deeper networks add capacity without much data to support it. For Credit, the one-hot feature representation is sparse and the tree-based models fit the rule-like structure more naturally. Naive Bayes was weaker than the best competing methods where it was tested, especially on Parkinson’s, which is reasonable because both the voice measurements and rice shape descriptors are likely correlated and violate the model’s conditional-independence assumption.

6 Extra Credit

[EC#1] Additional Algorithms on All Four Main Datasets

We also satisfy EC#1 by evaluating more than two algorithms on each of the four required datasets. Digits uses k-NN, Random Forest, and Neural Network; Parkinson’s uses k-NN, Gaussian Naive Bayes, Random Forest, and Neural Network; Rice uses k-NN, Decision Tree, and Gaussian Naive Bayes; and Credit Approval uses Decision Tree, Random Forest, and Neural Network. The detailed hyperparameter sweeps and best results for these additional algorithms are integrated directly into Sections 3.1–3.4 and summarized in the main results table.


[EC#2] Image Classification on Olivetti Faces

To satisfy the EC#2 criterion of selecting a new *challenging* dataset, we evaluated our k-NN (HW1) and Neural Network (HW4) implementations on the **Olivetti Faces** dataset — 400 grayscale 64x64 images of 40 distinct subjects (10 images per subject), loaded with a `scikit-learn` dataset fetcher. This is an image classification task with 40 classes, far more than any of the four main project datasets. The full notebook is at `notebooks/05_olivetti_ec2.ipynb`.


Random Forest was omitted because the from-scratch implementation is too slow with 4096 pixel features. k-NN and Neural Network are also strong face-recognition baselines.

Table 15: Olivetti Faces (EC#2): best result per algorithm, stratified 10-fold CV.

Algorithm	Best Hyperparameter	Accuracy	Macro F1
k-NN	$k = 1$	0.9575	0.9442
Neural Network	(128, 64), $\lambda = 0.0$	0.8350	0.8011



(a) k-NN: Accuracy and F1 vs k



(b) Neural Network: Accuracy and F1 by configuration

Figure 9: Olivetti Faces (EC#2): hyperparameter sweeps.




Figure 10: Olivetti Faces (EC#2): NN learning curve.

The Olivetti results support the extra-credit motivation: this is a small but difficult 40-class image-classification task. k-NN achieved the best result, with $k = 1$ reaching 0.9575 accuracy and 0.9442 macro F1. Performance decreased as k increased, which suggests that each subject's nearest images are highly informative but larger neighborhoods begin mixing other identities. The Neural Network also learned a useful classifier, with its best architecture reaching 0.8350 accuracy and 0.8011 macro F1, but it was weaker than k-NN. This is expected because the dataset has only 400 images and 40 classes, so the network has limited data for learning a stable face representation from 4096 pixel-level inputs.